

Muhammad Umer Adeeb

Question 4: Kaggle Dataset

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv('/content/Heart.csv')
df.head()

{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 303,\n  \"fields\": [\n    {\n      \"column\": \"Unnamed: 0\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 87,\n        \"min\": 1,\n        \"max\": 303,\n        \"num_unique_values\": 303,\n        \"samples\": [\n          180,\n          229,\n          112\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 9,\n        \"min\": 29,\n        \"max\": 77,\n        \"num_unique_values\": 41,\n        \"samples\": [\n          61,\n          64,\n          44\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Sex\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"ChestPain\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 4,\n        \"samples\": [\n          \"asymptomatic\",\n          \"nontypical\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"RestBP\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 17,\n        \"min\": 94,\n        \"max\": 200,\n        \"num_unique_values\": 50,\n        \"samples\": [\n          124,\n          192\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Chol\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 51,\n        \"min\": 126,\n        \"max\": 564,\n        \"num_unique_values\": 152,\n        \"samples\": [\n          321,\n          187\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Fbs\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}}
```



```
2   Sex      303 non-null   int64
3   ChestPain 303 non-null   object
4   RestBP   303 non-null   int64
5   Chol     303 non-null   int64
6   Fbs      303 non-null   int64
7   RestECG  303 non-null   int64
8   MaxHR    303 non-null   int64
9   ExAng    303 non-null   int64
10  Oldpeak  303 non-null   float64
11  Slope    303 non-null   int64
12  Ca       299 non-null   float64
13  Thal     301 non-null   object
14  AHD      303 non-null   object
dtypes: float64(2), int64(10), object(3)
memory usage: 35.6+ KB
```

```
df.shape
```

```
(303, 15)
```

```
df.isnull().sum()
```

```
Unnamed: 0    0
Age           0
Sex           0
ChestPain     0
RestBP        0
Chol          0
Fbs           0
RestECG       0
MaxHR         0
ExAng         0
Oldpeak       0
Slope         0
Ca            4
Thal          2
AHD           0
dtype: int64
```

```
df.dtypes
```

```
Unnamed: 0    int64
Age           int64
Sex           int64
ChestPain     object
RestBP        int64
Chol          int64
Fbs           int64
RestECG       int64
MaxHR         int64
ExAng         int64
```

```
Oldpeak      float64
Slope        int64
Ca           float64
Thal         object
AHD          object
dtype: object
```

```
for col in df.columns:
    zero_count = (df[col] == 0).sum()
    print(f"Number of zeros in {col}: {zero_count}")
```

```
Number of zeros in Unnamed: 0: 0
Number of zeros in Age: 0
Number of zeros in Sex: 97
Number of zeros in ChestPain: 0
Number of zeros in RestBP: 0
Number of zeros in Chol: 0
Number of zeros in Fbs: 258
Number of zeros in RestECG: 151
Number of zeros in MaxHR: 0
Number of zeros in ExAng: 204
Number of zeros in Oldpeak: 99
Number of zeros in Slope: 0
Number of zeros in Ca: 176
Number of zeros in Thal: 0
Number of zeros in AHD: 0
```

```
df['Age'].mean()
```

```
54.43894389438944
```

```
df.columns
```

```
Index(['Unnamed: 0', 'Age', 'Sex', 'ChestPain', 'RestBP', 'Chol',
       'Fbs',
       'RestECG', 'MaxHR', 'ExAng', 'Oldpeak', 'Slope', 'Ca', 'Thal',
       'AHD'],
      dtype='object')
```

```
#Now extract only Age, Sex, ChestPain, RestBP, Chol.
```

```
df.drop(columns=['Unnamed: 0', 'Fbs', 'RestECG', 'MaxHR', 'ExAng',
                 'Oldpeak', 'Slope', 'Ca', 'Thal'], inplace=True)
df.head()
```

```
{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 303,\n  \"fields\": [\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 9,\n        \"min\": 29,\n        \"max\": 77,\n        \"num_unique_values\": 41,\n        \"samples\": [\n          61,\n          64,\n          44\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Sex\",\n      \"properties\": {\n
```

```

{"dtype\\": \\\"number\\",\\n          \\\"std\\\": 0,\\n          \\\"min\\\": 0,\\n
\\\"max\\\": 1,\\n          \\\"num_unique_values\\\": 2,\\n          \\\"samples\\\":
[\\n          0,\\n          1\\n          ],\\n          \\\"semantic_type\\\":
\\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n          }\\n          },\\n          {\\n
\\\"column\\\": \\\"ChestPain\\\",\\n          \\\"properties\\\": {\\n
\\\"dtype\\\": \\\"category\\\",\\n          \\\"num_unique_values\\\": 4,\\n
\\\"samples\\\": [\\n          \\\"asymptomatic\\\",\\n          \\\"nontypical\\\"\\n
          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n          }\\n          },\\n          {\\n          \\\"column\\\":
\\\"RestBP\\\",\\n          \\\"properties\\\": {\\n          \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 17,\\n          \\\"min\\\": 94,\\n          \\\"max\\\": 200,\\n
\\\"num_unique_values\\\": 50,\\n          \\\"samples\\\": [\\n          124,\\n
192\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n          }\\n          },\\n          {\\n          \\\"column\\\":
\\\"Chol\\\",\\n          \\\"properties\\\": {\\n          \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 51,\\n          \\\"min\\\": 126,\\n          \\\"max\\\": 564,\\n
\\\"num_unique_values\\\": 152,\\n          \\\"samples\\\": [\\n          321,\\n
187\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n          }\\n          },\\n          {\\n          \\\"column\\\":
\\\"AHD\\\",\\n          \\\"properties\\\": {\\n          \\\"dtype\\\": \\\"category\\\",\\n
\\\"num_unique_values\\\": 2,\\n          \\\"samples\\\": [\\n          \\\"Yes\\\",\\n
          \\\"No\\\"\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n          }\\n          }\\n          ]\\n
n}\\", \"type\": \"dataframe\", \"variable_name\": \"df\"}

```

```

x = df.drop(columns=['AHD'])
y= df['AHD']
display(x.head())
display(y.head())

```

```

{"summary": "{\\n  \\\"name\\\": \\\"display(y\\\",\\n  \\\"rows\\\": 5,\\n
\\\"fields\\\": [\\n    {\\n      \\\"column\\\": \\\"Age\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\":
14,\\n        \\\"min\\\": 37,\\n        \\\"max\\\": 67,\\n
\\\"num_unique_values\\\": 4,\\n        \\\"samples\\\": [\\n        67,\\n
41,\\n        63\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n        }\\n        },\\n        {\\n        \\\"column\\\":
\\\"Sex\\\",\\n        \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 0,\\n        \\\"min\\\": 0,\\n        \\\"max\\\": 1,\\n
\\\"num_unique_values\\\": 2,\\n        \\\"samples\\\": [\\n        0,\\n
1\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n        }\\n        },\\n        {\\n        \\\"column\\\":
\\\"ChestPain\\\",\\n        \\\"properties\\\": {\\n        \\\"dtype\\\":
\\\"string\\\",\\n        \\\"num_unique_values\\\": 4,\\n        \\\"samples\\\":
[\\n        \\\"asymptomatic\\\",\\n        \\\"nontypical\\\"\\n        ],\\n
\\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n        }\\n
        },\\n        {\\n        \\\"column\\\": \\\"RestBP\\\",\\n        \\\"properties\\\":
{\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\": 15,\\n
\\\"min\\\": 120,\\n        \\\"max\\\": 160,\\n        \\\"num_unique_values\\\":
4,\\n        \\\"samples\\\": [\\n        160,\\n        130\\n        ],\\n

```

```

n      \ "semantic_type\": \ "\",\n      \ "description\": \ "\n
}\n    },\n    {\n      \ "column\": \ "Chol\",\n      \ "properties\":
{\n      \ "dtype\": \ "number\",\n      \ "std\": 30,\n
\ "min\": 204,\n      \ "max\": 286,\n      \ "num_unique_values\":
5,\n      \ "samples\": [\n      286,\n      204\n      ],\n
n      \ "semantic_type\": \ "\",\n      \ "description\": \ "\n
}\n    }\n  ]\n}", "type": "dataframe"}

```

```

0      No
1      Yes
2      Yes
3      No
4      No

```

Name: AHD, dtype: object

```

y.replace({'Yes': 1, 'No': 0}, inplace=True)
display(y.head())

```

<ipython-input-12-a8acbb79871c>:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`

```

y.replace({'Yes': 1, 'No': 0}, inplace=True)

```

```

0      0
1      1
2      1
3      0
4      0

```

Name: AHD, dtype: int64

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,
test_size=0.25, random_state=90)
print(x_train.shape)
display(x_train.head())
print(x_test.shape)
display(x_test.head())
print(y_train.shape)
display(y_train.head())
print(y_test.shape)
display(y_test.head())

```

(227, 5)

```

{"summary": "{\n  \ "name\": \ "display(y_test\",\n  \ "rows\": 5,\n
\ "fields\": [\n    {\n      \ "column\": \ "Age\",\n
\ "properties\": {\n      \ "dtype\": \ "number\",\n      \ "std\":
8,\n      \ "min\": 40,\n      \ "max\": 62,\n
\ "num_unique_values\": 5,\n      \ "samples\": [\n      59,\n

```

```

55,\n          57\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n          },\n          {\n          \"column\":\n          \"Sex\",\n          \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 0,\n          \"min\": 0,\n          \"max\": 1,\n          \"num_unique_values\": 2,\n          \"samples\": [\n          1,\n          0\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n          },\n          {\n          \"column\":\n          \"ChestPain\",\n          \"properties\": {\n          \"dtype\":\n          \"string\",\n          \"num_unique_values\": 3,\n          \"samples\":\n          [\n          \"asymptomatic\",\n          \"typical\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n          },\n          {\n          \"column\": \"RestBP\",\n          \"properties\":\n          {\n          \"dtype\": \"number\",\n          \"std\": 19,\n          \"min\": 132,\n          \"max\": 178,\n          \"num_unique_values\":\n          4,\n          \"samples\": [\n          178,\n          140\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n          },\n          {\n          \"column\": \"Chol\",\n          \"properties\":\n          {\n          \"dtype\": \"number\",\n          \"std\": 57,\n          \"min\": 199,\n          \"max\": 342,\n          \"num_unique_values\":\n          5,\n          \"samples\": [\n          270,\n          342\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n          }\n          }\n          ],\n          \"type\": \"dataframe\"}

```

(76, 5)

```

{"summary": "{\n  \"name\": \"display(y_test\",\n  \"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 3,\n        \"min\": 53,\n        \"max\": 61,\n        \"num_unique_values\": 4,\n        \"samples\": [\n          53,\n          59,\n          56\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\":\n        \"Sex\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 0,\n          \"min\": 1,\n          \"max\": 1,\n          \"num_unique_values\": 1,\n          \"samples\": [\n          1\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"ChestPain\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 2,\n          \"samples\": [\n          \"asymptomatic\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\":\n          \"RestBP\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 11,\n            \"min\": 110,\n            \"max\": 138,\n            \"num_unique_values\": 4,\n            \"samples\": [\n            123\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"Chol\",\n            \"properties\":\n            {\n              \"dtype\": \"number\",\n              \"std\": 46,\n              \"min\": 166,\n              \"max\": 282,\n              \"num_unique_values\":\n              5,\n              \"samples\": [\n              282\n              ],\n

```

```
\ "semantic_type\": \ "\",\n          \ "description\": \ "\",\n          }\n        ]\n}", "type": "dataframe"}
```

```
(227,)
```

```
209    1
183    0
165    0
41     0
291    0
```

```
Name: AHD, dtype: int64
```

```
(76,)
```

```
<bound method NDFrame.head of 288    0
```

```
223    1
264    1
96     1
219    0
```

```
..
55     1
79     1
217    0
287    0
188    1
```

```
Name: AHD, Length: 76, dtype: int64>
```

```
x_train['ChestPain'].unique()
```

```
array(['asymptomatic', 'typical', 'nontypical', 'nonanginal'],
      dtype=object)
```

```
from sklearn.preprocessing import OrdinalEncoder
```

```
ordinal_encoder = OrdinalEncoder(categories=[['typical',
'asymptomatic', 'nonanginal', 'nontypical']])
```

```
x_train['ChestPain'] =
```

```
ordinal_encoder.fit_transform(x_train[['ChestPain']])
```

```
x_test['ChestPain'] = ordinal_encoder.transform(x_test[['ChestPain']])
```

```
print(x_train.head())
```

```
print(x_test.head())
```

	Age	Sex	ChestPain	RestBP	Chol
209	62	0	1.0	150	244
183	59	1	0.0	178	270
165	57	1	1.0	132	207
41	40	1	0.0	140	199
291	55	0	3.0	132	342
	Age	Sex	ChestPain	RestBP	Chol
288	56	1	3.0	130	221
223	53	1	1.0	123	282

264	61	1	1.0	138	166
96	59	1	1.0	110	239
219	59	1	1.0	138	271

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
x_train[['Age', 'RestBP', 'Chol']] =
scaler.fit_transform(x_train[['Age', 'RestBP', 'Chol']])
x_test[['Age', 'RestBP', 'Chol']] = scaler.transform(x_test[['Age',
'RestBP', 'Chol']])
```

```
print(x_train.head())
print(x_test.head())
```

	Age	Sex	ChestPain	RestBP	Chol
209	0.834084	0	1.0	1.014683	-0.025148
183	0.508203	1	0.0	2.627015	0.469594
165	0.290948	1	1.0	-0.021816	-0.729203
41	-1.555713	1	0.0	0.438850	-0.881431
291	0.073694	0	3.0	-0.021816	1.839647
	Age	Sex	ChestPain	RestBP	Chol
288	0.182321	1	3.0	-0.136982	-0.462804
223	-0.143560	1	1.0	-0.540065	0.697936
264	0.725457	1	1.0	0.323684	-1.509372
96	0.508203	1	1.0	-1.288648	-0.120290
219	0.508203	1	1.0	0.323684	0.488622

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train, y_train)
```

```
LogisticRegression()
```

```
y_pred = model.predict(x_test)
y_pred
```

```
array([0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1,
0,
      1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1,
      0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0,
0,
      1, 0, 0, 0, 1, 1, 1, 0, 0, 0])
```

```
pd.DataFrame({'y_test':y_test, 'y_predict':y_pred})
```

```
{"summary":{"\n  \"name\": \"pd\",\n  \"rows\": 76,\n  \"fields\": [\n    {\n      \"column\": \"y_test\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\":
```

```
[\\n          1,\\n          0\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n          }\\n          },\\n          {\\n          \\\"column\\\": \\\"y_predict\\\",\\n          \\\"properties\\\": {\\n          \\\"dtype\\\": \\\"number\\\",\\n          \\\"std\\\": 0,\\n          \\\"min\\\": 0,\\n          \\\"max\\\": 1,\\n          \\\"num_unique_values\\\": 2,\\n          \\\"samples\\\": \\n          1,\\n          0\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n          }\\n          }\\n          ]\\n          }\\\", \"type\": \"dataframe\"}
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
print('Accuracy Score: ', accuracy_score(y_test, y_pred))
print('Precision Score: ', precision_score(y_test, y_pred))
print('Recall Score: ', recall_score(y_test, y_pred))
print('F1 Score: ', f1_score(y_test, y_pred))
print('Confusion Matrix:\\n ', confusion_matrix(y_test, y_pred))
print('Classification Report:\\n ', classification_report(y_test, y_pred))
```

Accuracy Score: 0.75

Precision Score: 0.7428571428571429

Recall Score: 0.7222222222222222

F1 Score: 0.7323943661971831

Confusion Matrix:

```
[[31  9]
```

```
[10 26]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.76	0.78	0.77	40
1	0.74	0.72	0.73	36
accuracy			0.75	76
macro avg	0.75	0.75	0.75	76
weighted avg	0.75	0.75	0.75	76

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

